

```
from google.colab import files
uploaded = files.upload()
```



Choose Files WA\_Fn-Us...-Attrition.csv

- **WA\_Fn-UseC\_-HR-Employee-Attrition.csv**(text/csv) - 227977 bytes, last modified: 9/20/2019 - 100% done

Saving WA\_Fn-UseC\_-HR-Employee-Attrition.csv to WA\_Fn-UseC\_-HR-Employee-Attrition.csv

```
import pandas as pd
```

```
# Load the dataset
```

```
df = pd.read_csv("WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

```
# View shape and first few rows
```

```
print("Dataset shape:", df.shape)
```

```
df.head()
```



Dataset shape: (1470, 35)

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	...	RelationshipSatisfaction	Standard
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	...	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	...	4	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	...	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	...	3	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	...	4	

5 rows × 35 columns

```
# Overview of dataset
```

```
df.info()
```

```
# Check for missing values
```

```
df.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1470 non-null   int64
1   Attrition                            1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                          1470 non-null   object
5   DistanceFromHome                    1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                      1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EmployeeNumber                      1470 non-null   int64
10  EnvironmentSatisfaction              1470 non-null   int64
11  Gender                              1470 non-null   object
12  HourlyRate                          1470 non-null   int64
13  JobInvolvement                      1470 non-null   int64
14  JobLevel                            1470 non-null   int64
15  JobRole                             1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                      1470 non-null   object
18  MonthlyIncome                      1470 non-null   int64
19  MonthlyRate                        1470 non-null   int64
20  NumCompaniesWorked                 1470 non-null   int64
21  Over18                             1470 non-null   object
22  OverTime                           1470 non-null   object
23  PercentSalaryHike                  1470 non-null   int64
24  PerformanceRating                  1470 non-null   int64
25  RelationshipSatisfaction            1470 non-null   int64
26  StandardHours                     1470 non-null   int64
27  StockOptionLevel                   1470 non-null   int64
28  TotalWorkingYears                  1470 non-null   int64
29  TrainingTimesLastYear              1470 non-null   int64
30  WorkLifeBalance                    1470 non-null   int64
31  YearsAtCompany                     1470 non-null   int64
32  YearsInCurrentRole                 1470 non-null   int64
33  YearsSinceLastPromotion            1470 non-null   int64
34  YearsWithCurrManager               1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

	0
Age	0
Attrition	0
BusinessTravel	0
DailyRate	0
Department	0
DistanceFromHome	0

Education	0
EducationField	0
EmployeeCount	0
EmployeeNumber	0
EnvironmentSatisfaction	0
Gender	0
HourlyRate	0
JobInvolvement	0
JobLevel	0
JobRole	0
JobSatisfaction	0
MaritalStatus	0
MonthlyIncome	0
MonthlyRate	0
NumCompaniesWorked	0
Over18	0
OverTime	0
PercentSalaryHike	0
PerformanceRating	0
RelationshipSatisfaction	0
StandardHours	0
StockOptionLevel	0
TotalWorkingYears	0
TrainingTimesLastYear	0
WorkLifeBalance	0
YearsAtCompany	0
YearsInCurrentRole	0
YearsSinceLastPromotion	0
YearsWithCurrManager	0

dtype: int64



```
# Label encode the target variable
df['Attrition'] = df['Attrition'].map({'Yes': 1, 'No': 0})

# Get all object (text) columns
categorical_cols = df.select_dtypes(include='object').columns
print("Categorical columns:", list(categorical_cols))

🔗 Categorical columns: ['BusinessTravel', 'Department', 'EducationField', 'Gender', 'JobRole', 'MaritalStatus', 'Over18', 'OverTime']

# One-hot encode all remaining categorical columns
df_encoded = pd.get_dummies(df, drop_first=True)

# Confirm all columns are now numeric
df_encoded.info()

# Preview the encoded dataset
df_encoded.head()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 48 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   Age                                    1470 non-null   int64
 1   Attrition                             1470 non-null   int64
 2   DailyRate                             1470 non-null   int64
 3   DistanceFromHome                      1470 non-null   int64
 4   Education                             1470 non-null   int64
 5   EmployeeCount                         1470 non-null   int64
 6   EmployeeNumber                        1470 non-null   int64
 7   EnvironmentSatisfaction               1470 non-null   int64
 8   HourlyRate                            1470 non-null   int64
 9   JobInvolvement                       1470 non-null   int64
10  JobLevel                              1470 non-null   int64
11  JobSatisfaction                       1470 non-null   int64
12  MonthlyIncome                        1470 non-null   int64
13  MonthlyRate                           1470 non-null   int64
14  NumCompaniesWorked                   1470 non-null   int64
15  PercentSalaryHike                    1470 non-null   int64
16  PerformanceRating                    1470 non-null   int64
17  RelationshipSatisfaction              1470 non-null   int64
18  StandardHours                        1470 non-null   int64
19  StockOptionLevel                     1470 non-null   int64
20  TotalWorkingYears                    1470 non-null   int64
21  TrainingTimesLastYear                1470 non-null   int64
22  WorkLifeBalance                      1470 non-null   int64
23  YearsAtCompany                       1470 non-null   int64
24  YearsInCurrentRole                   1470 non-null   int64
25  YearsSinceLastPromotion               1470 non-null   int64
26  YearsWithCurrManager                 1470 non-null   int64
27  BusinessTravel_Travel_Frequently     1470 non-null   bool
28  BusinessTravel_Travel_Rarely         1470 non-null   bool
29  Department_Research & Development    1470 non-null   bool
30  Department_Sales                     1470 non-null   bool
31  EducationField_Life Sciences          1470 non-null   bool
32  EducationField_Marketing              1470 non-null   bool
33  EducationField_Medical                1470 non-null   bool
34  EducationField_Other                  1470 non-null   bool
35  EducationField_Technical Degree       1470 non-null   bool
36  Gender_Male                           1470 non-null   bool
37  JobRole_Human Resources                1470 non-null   bool
38  JobRole_Laboratory Technician         1470 non-null   bool
39  JobRole_Manager                       1470 non-null   bool
40  JobRole_Manufacturing Director        1470 non-null   bool
41  JobRole_Research Director             1470 non-null   bool
42  JobRole_Research Scientist            1470 non-null   bool
43  JobRole_Sales Executive                1470 non-null   bool
44  JobRole_Sales Representative           1470 non-null   bool
45  MaritalStatus_Married                 1470 non-null   bool
46  MaritalStatus_Single                  1470 non-null   bool
47  OverTime_Yes                          1470 non-null   bool
dtypes: bool(21), int64(27)
memory usage: 340.4 KB

```

	Age	Attrition	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	HourlyRate	JobInvolvement	...	JobRole_Laboratory Technician	JobR
0	41	1	1102	1	2	1	1	2	94	3	...	False	
1	49	0	279	8	1	1	2	3	61	2	...	False	
2	37	1	1373	2	2	1	4	4	92	2	...	True	
3	33	0	1392	3	4	1	5	4	56	3	...	False	
4	27	0	591	2	1	1	7	1	40	3	...	True	

5 rows × 48 columns

```
# Check correlation with 'Attrition'
corr = df_encoded.corr()['Attrition'].sort_values(ascending=False)
print(corr)
```

```
Attrition      1.000000
OverTime_Yes   0.246118
MaritalStatus_Single 0.175419
JobRole_Sales Representative 0.157234
BusinessTravel_Travel_Frequently 0.115143
JobRole_Laboratory Technician 0.098290
Department_Sales 0.080855
DistanceFromHome 0.077924
EducationField_Technical Degree 0.069355
EducationField_Marketing 0.055781
NumCompaniesWorked 0.043494
JobRole_Human Resources 0.036215
Gender_Male     0.029453
JobRole_Sales Executive 0.019774
MonthlyRate     0.015170
PerformanceRating 0.002889
JobRole_Research Scientist -0.000360
HourlyRate      -0.006846
EmployeeNumber  -0.010577
PercentSalaryHike -0.013478
EducationField_Other -0.017898
Education       -0.031373
EducationField_Life Sciences -0.032703
YearsSinceLastPromotion -0.033019
RelationshipSatisfaction -0.045872
EducationField_Medical -0.046999
BusinessTravel_Travel_Rarely -0.049538
DailyRate       -0.056652
TrainingTimesLastYear -0.059478
WorkLifeBalance -0.063939
```


```
JobRole_Manufacturing Director    -0.082994
JobRole_Manager                   -0.083316
Department_Research & Development -0.085293
JobRole_Research Director         -0.088870
MaritalStatus_Married             -0.090984
EnvironmentSatisfaction            -0.103369
JobSatisfaction                   -0.103481
JobInvolvement                    -0.130016
YearsAtCompany                    -0.134392
StockOptionLevel                  -0.137145
YearsWithCurrManager              -0.156199
Age                               -0.159205
MonthlyIncome                     -0.159840
YearsInCurrentRole                 -0.160545
JobLevel                           -0.169105
TotalWorkingYears                 -0.171063
EmployeeCount                     NaN
StandardHours                     NaN
Name: Attrition, dtype: float64
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Drop Attrition itself and get top correlations
top_corr_features = corr.drop('Attrition').head(10)
```

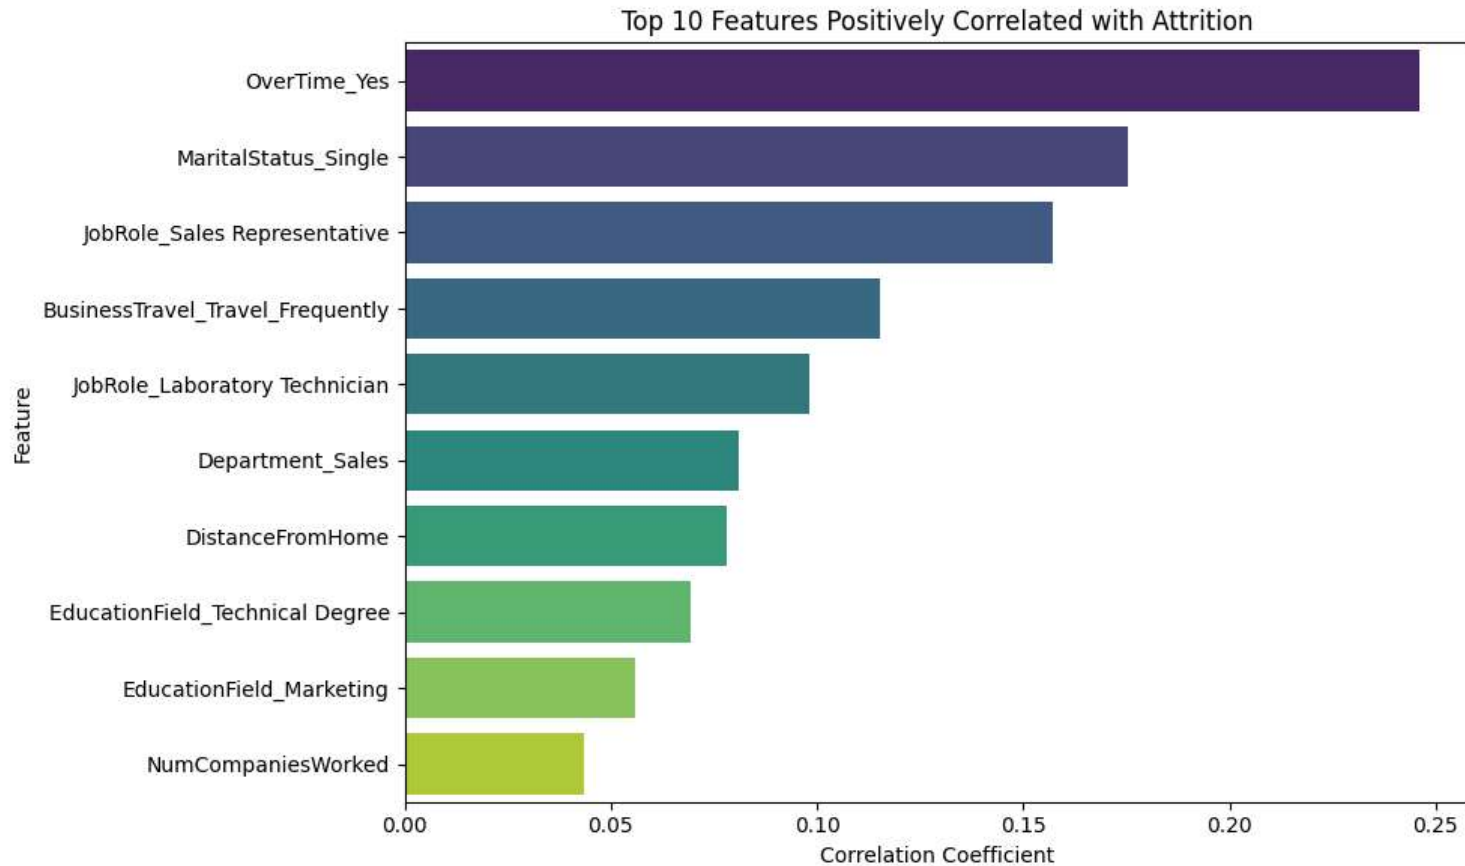
```
plt.figure(figsize=(10, 6))
sns.barplot(x=top_corr_features.values, y=top_corr_features.index, palette="viridis")
plt.title("Top 10 Features Positively Correlated with Attrition")
plt.xlabel("Correlation Coefficient")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```




 <ipython-input-9-9e38d110ffad>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=top_corr_features.values, y=top_corr_features.index, palette="viridis")
```



```
# Example: MonthlyIncome vs Attrition
sns.boxplot(x='Attrition', y='MonthlyIncome', data=df)
plt.title("Monthly Income vs Attrition")
plt.show()
```



```
# Define X and y
X = df_encoded.drop('Attrition', axis=1)
y = df_encoded['Attrition']

from sklearn.model_selection import train_test_split

# 80/20 split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Initialize models
dt = DecisionTreeClassifier(random_state=42)
rf = RandomForestClassifier(random_state=42)

# Train models
dt.fit(X_train, y_train)
rf.fit(X_train, y_train)
```



RandomForestClassifier



RandomForestClassifier(random\_state=42)

```
# Predict
dt_pred = dt.predict(X_test)
rf_pred = rf.predict(X_test)

# Define a function for evaluation
def evaluate_model(name, y_true, y_pred):
    print(f"\n{name} Results:")
    print("Accuracy:", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred))
    print("Recall:", recall_score(y_true, y_pred))
    print("F1 Score:", f1_score(y_true, y_pred))

# Evaluate both
evaluate_model("Decision Tree", y_test, dt_pred)
evaluate_model("Random Forest", y_test, rf_pred)
```



```
Decision Tree Results:
Accuracy: 0.7755102040816326
Precision: 0.17073170731707318
Recall: 0.1794871794871795
F1 Score: 0.175

Random Forest Results:
Accuracy: 0.8775510204081632
Precision: 0.8
Recall: 0.10256410256410256
F1 Score: 0.18181818181818182
```


```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Get feature importances
importances = rf.feature_importances_
features = X.columns

# Create a DataFrame
feat_df = pd.DataFrame({'Feature': features, 'Importance': importances})
feat_df = feat_df.sort_values(by='Importance', ascending=False)

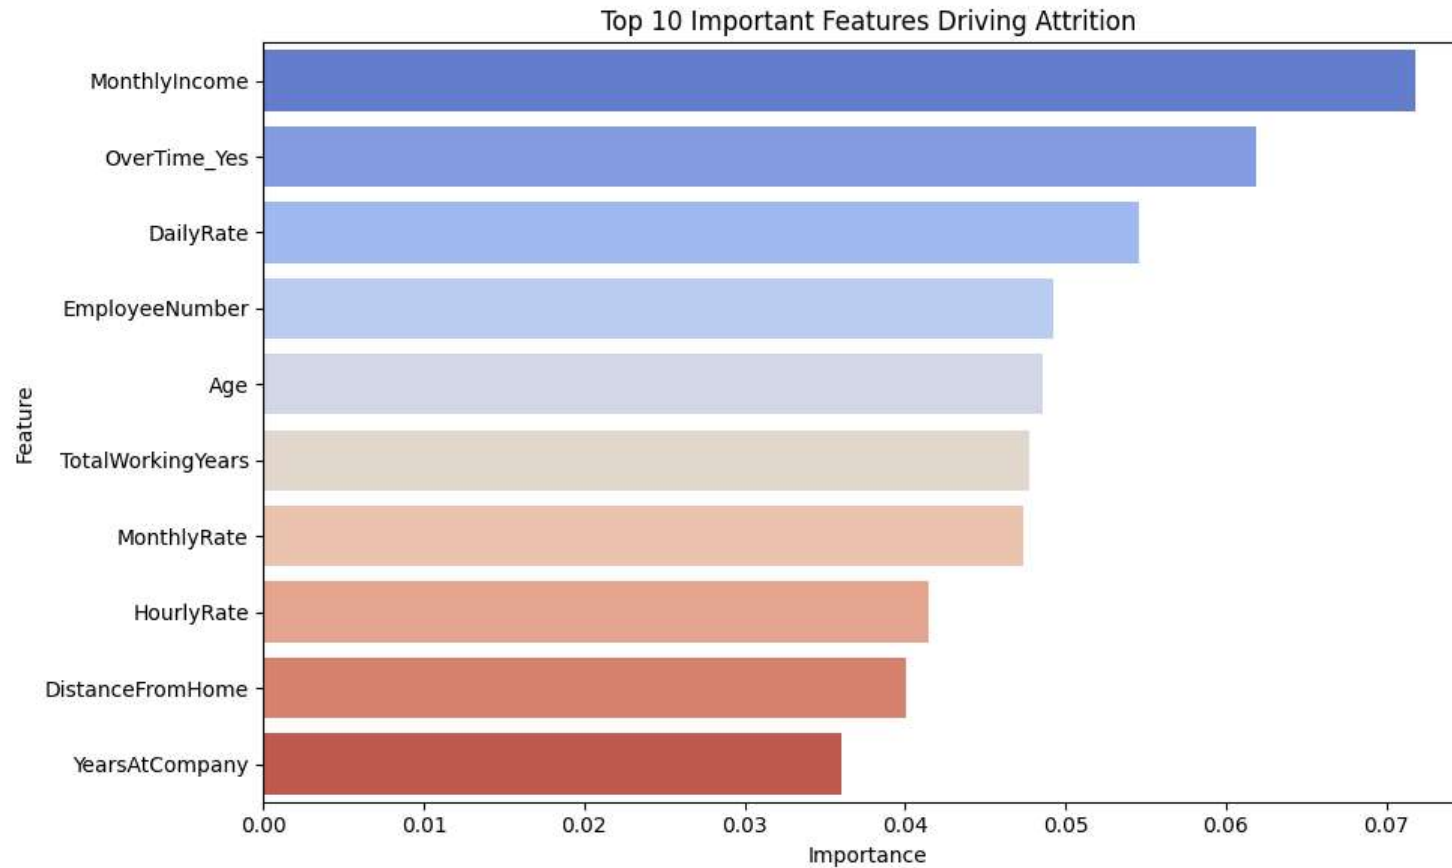
# Plot top 10 important features
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=feat_df.head(10), palette="coolwarm")
plt.title('Top 10 Important Features Driving Attrition')
```

```
plt.tight_layout()
plt.show()
```

 <ipython-input-16-d6fcb16262f4>:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Importance', y='Feature', data=feat_df.head(10), palette="coolwarm")
```



```
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
import xgboost as xgb
from sklearn.ensemble import VotingClassifier
```

```
# Initialize models
lr = LogisticRegression(max_iter=1000)
knn = KNeighborsClassifier()
```


```
gb = GradientBoostingClassifier()
xgb_model = xgb.XGBClassifier(use_label_encoder=False, eval_metric='logloss')
```

```
# Voting classifier (ensemble of top models)
voting_clf = VotingClassifier(estimators=[
    ('lr', lr), ('knn', knn), ('rf', rf), ('gb', gb), ('xgb', xgb_model)
], voting='soft')
```

```
models = {
    "Logistic Regression": lr,
    "K-Nearest Neighbors": knn,
    "Gradient Boosting": gb,
    "XGBoost": xgb_model,
    "Voting Classifier": voting_clf
}
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f"\n{name}")
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("Precision:", precision_score(y_test, y_pred))
    print("Recall:", recall_score(y_test, y_pred))
    print("F1 Score:", f1_score(y_test, y_pred))
```

 /usr/local/lib/python3.11/dist-packages/sklearn/linear\_model/\_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Logistic Regression

Accuracy: 0.8673469387755102

Precision: 0.5

Recall: 0.1794871794871795

F1 Score: 0.2641509433962264

K-Nearest Neighbors

Accuracy: 0.8537414965986394

Precision: 0.35714285714285715

Recall: 0.1282051282051282

F1 Score: 0.18867924528301888

Gradient Boosting

Accuracy: 0.8775510204081632

Precision: 0.6

Recall: 0.23076923076923078